



# CSC4005 – Distributed and Parallel Computing

Prof. Yeh-Ching Chung

School of Data Science

Chinese University of Hong Kong, Shenzhen





# Outline

- Introduction to Parallel Computers
- Message Passing Computing and Programming
- Multithreaded Programming
- CUDA Programming
- OpenMP Programming
- **Embarrassingly Parallel Computations**
- Partitioning and Divide-and-Conquer Strategies
- Pipelined Computations
- Synchronous Computations
- Load Balancing and Termination Detection
- Sorting Algorithms





# Embarrassingly Parallel Computations (1)

A computation that can be divided into a number of completely independent parts, each of which can be executed by a separate processor.




Figure 3.1 Disconnected computational graph (embarrassingly parallel problem).





## Embarrassingly Parallel Computations (2)




Figure 3.2 Practical embarrassingly parallel computational graph with dynamic process creation and the master-slave approach.





# Embarrassingly Parallel Examples (1)

## Low level image operations:

### (a) Shifting

Object shifted by  $\Delta x$  in the  $x$ -dimension and  $\Delta y$  in the  $y$ -dimension:

$$x' = x + \Delta x$$

$$y' = y + \Delta y$$

where  $x$  and  $y$  are the original and  $x'$  and  $y'$  are the new coordinates.

### (b) Scaling

Object scaled by a factor  $S_x$  in the  $x$ -direction and  $S_y$  in the  $y$ -direction:

$$x' = xS_x$$

$$y' = yS_y$$

### (c) Rotation

Object rotated through an angle  $\theta$  about the origin of the coordinate system:


$$x' = x \cos \theta + y \sin \theta$$

$$y' = -x \sin \theta + y \cos \theta$$





## Embarrassingly Parallel Examples (2)




(a) Square region for each process

Figure 3.3 Partitioning into regions for individual processes.





## Embarrassingly Parallel Examples (3)



(b) Row region for each process





# Pseudocode to Perform Image Shift (1)

Master

```
for (i = 0, row = 0; i < 48; i++, row = row + 10) /* for each process*/
    send(row, Pi);                                /* send row no.*/

for (i = 0; i < 480; i++)                          /* initialize temp */
    for (j = 0; j < 640; j++)
        temp_map[i][j] = 0;

for (i = 0; i < (640 * 480); i++) {                /* for each pixel */
    recv(oldrow, oldcol, newrow, newcol, PANY); /* accept new coords */
    if (!((newrow < 0) || (newrow >= 480) || (newcol < 0) || (newcol >= 640))
        temp_map[newrow][newcol] = map[oldrow][oldcol];
}

for (i = 0; i < 480; i++)                          /* update bitmap */
    for (j = 0; j < 640; j++)
        map[i][j] = temp_map[i][j];
```







## Embarrassingly Parallel Examples (2)

Slave

```
recv(row, Pmaster);          /* receive row no. */
for (oldrow = row; oldrow < (row + 10); oldrow++)
    for (oldcol = 0; oldcol < 640; oldcol++) { /* transform coords */
        newrow = oldrow + delta_x;          /* shift in x direction */
        newcol = oldcol + delta_y;          /* shift in y direction */
        send(oldrow,oldcol,newrow,newcol, Pmaster); /* coords to master */
    }
```





## Embarrassingly Parallel Examples (3)

### Analysis

#### Sequential

$$t_s = n^2 = O(n^2)$$

#### Parallel Communication

$$t_{\text{comm}} = t_{\text{startup}} + mt_{\text{data}}$$

$$t_{\text{comm}} = p(t_{\text{startup}} + 2t_{\text{data}}) + 4n^2(t_{\text{startup}} + t_{\text{data}}) = O(p + n^2)$$

#### Computation

$$t_{\text{comp}} = 2\left(\frac{n^2}{p}\right) = O(n^2/p)$$

#### Overall Execution Time

$$t_p = t_{\text{comp}} + t_{\text{comm}}$$

For constant  $p$ , this is  $O(n^2)$ . However, the constant hidden in the communication part far exceeds those constants in the computation in most practical situations.





# Mandelbrot Set Computation (1)

Set of points in a complex plane that are quasi-stable (will increase and decrease, but not exceed some limit) when computed by iterating the function

$$z_{k+1} = z_k^2 + c$$

where  $z_{k+1}$  is the  $(k+1)$ th iteration of the complex number  $z = a + bi$  and  $c$  is a complex number giving the position of the point in the complex plane.

The initial value for  $z$  is zero.

Iterations continued until magnitude of  $z$  is greater than 2 or number of iterations reaches arbitrary limit. Magnitude of  $z$  is the length of the vector given by

$$z_{\text{length}} = \sqrt{a^2 + b^2}$$





## Mandelbrot Set Computation (2)

Computing the complex function,  $z_{k+1} = z_k^2 + c$ , is simplified by recognizing that

$$z^2 = a^2 + 2abi + bi^2 = a^2 - b^2 + 2abi$$

or a real part that is  $a^2 - b^2$  and an imaginary part that is  $2ab$ .

The next iteration values can be produced by computing:

$$z_{\text{real}} = z_{\text{real}}^2 - z_{\text{imag}}^2 + c_{\text{real}}$$

$$z_{\text{imag}} = 2z_{\text{real}}z_{\text{imag}} + c_{\text{imag}}$$





## Mandelbrot Set Computation (3)

**Seq. Routine computing value of one pt, returning no of iterations**

```
structure complex {
    float real;
    float imag;
};

int cal_pixel(complex c)
{
    int count, max;
    complex z;
    float temp, lengthsq;
    max = 256;
    z.real = 0; z.imag = 0;
    count = 0;                                /* number of iterations */
    do {
        temp = z.real * z.real - z.imag * z.imag + c.real;
        z.imag = 2 * z.real * z.imag + c.imag;
        z.real = temp;
        lengthsq = z.real * z.real + z.imag * z.imag;
        count++;
    } while ((lengthsq < 4.0) && (count < max));
    return count;
}
```





# Mandelbrot Set Computation (4)

## Scaling Coordinate System

For computational efficiency, let

```
scale_real = (real_max - real_min)/disp_width;  
scale_imag = (imag_max - imag_min)/disp_height;
```

Including scaling, the code could be of the form

```
for (x = 0; x < disp_width; x++) /* screen coordinates x and y */  
  for (y = 0; y < disp_height; y++) {  
    c.real = real_min + ((float) x * scale_real);  
    c.imag = imag_min + ((float) y * scale_imag);  
    color = cal_pixel(c);  
    display(x, y, color);  
  }
```

where `display()` is a routine to display the pixel (x, y) at the computed color.





## Mandelbrot Set Computation (5)




Figure 3.4 Mandelbrot set.





# Parallelization of Mandelbrot Computation (1)

## Static Task Assignment

Master

```
for (i = 0, row = 0; i < 48; i++, row = row + 10) /* for each process*/
    send(&row, Pi);                               /* send row no.*/
for (i = 0; i < (480 * 640); i++) { /* from processes, any order */
    recv(&c, &color, PANY); /* receive coordinates/colors */
    display(c, color);      /* display pixel on screen */
}
```

Slave (process i)

```
recv(&row, Pmaster); /* receive row no. */
for (x = 0; x < disp_width; x++) /* screen coordinates x and y */
    for (y = row; y < (row + 10); y++) {
        c.real = min_real + ((float) x * scale_real);
        c.imag = min_imag + ((float) y * scale_imag);
        color = cal_pixel(c);
        send(&c, &color, Pmaster); /* send coords, color to master */
    }
```







# Parallelization of Mandelbrot Computation (2)

## Dynamic Task Assignment Work Pool/Processor Farms

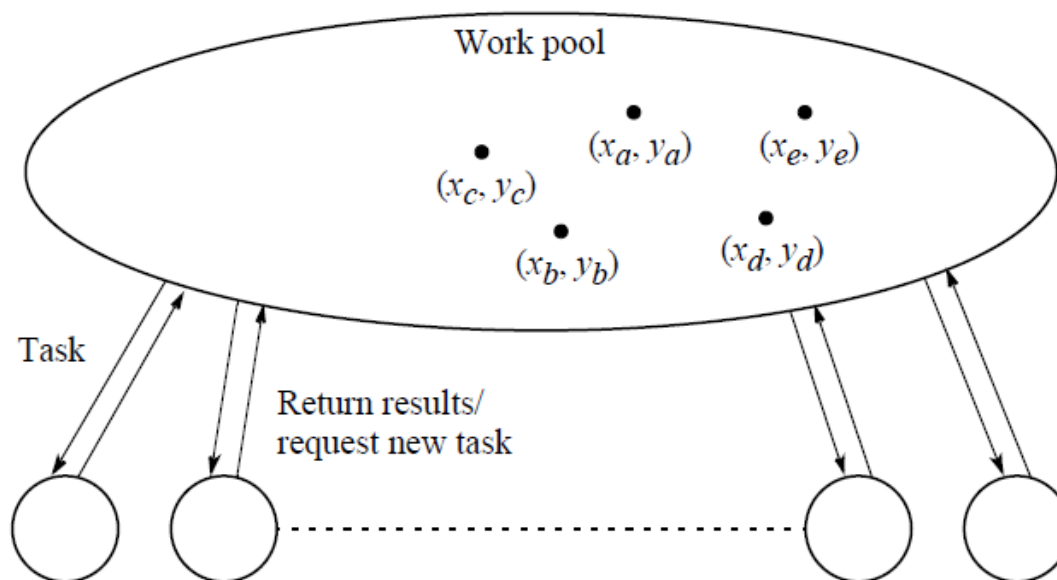


Figure 3.5 Work pool approach.





# Parallelization of Mandelbrot Computation (3)

## Coding for Work Pool Approach

Master

```
count = 0;                                /* counter for termination*/
row = 0;                                  /* row being sent */
for (k = 0; k < procno; k++) { /* assuming procno < disp_height */
    send(&row, Pk, data_tag); /* send initial row to process */
    count++;                    /* count rows sent */
    row++;                      /* next row */
}

do {
    recv (&slave, &r, color, PANY, result_tag);
    count--; /* reduce count as rows received */
    if (row < disp_height) {
        send (&row, Pslave, data_tag); /* send next row */
        row++; /* next row */
        count++;
    } else
        send (&row, Pslave, terminator_tag); /* terminate */
    rows_rcv++;
    display (r, color); /* display row */
} while (count > 0);
```





## Parallelization of Mandelbrot Computation (4)

Slave

```
recv(y, P_master, ANYTAG, source_tag); /* receive 1st row to compute */
while (source_tag == data_tag) {
    c.imag = imag_min + ((float) y * scale_imag);
    for (x = 0; x < disp_width; x++) { /* compute row colors */
        c.real = real_min + ((float) x * scale_real);
        color[x] = cal_pixel(c);
    }
    send(&i, &y, color, P_master, result_tag); /* row colors to master */
    recv(y, P_master, source_tag); /* receive next row */
};
```





## Parallelization of Mandelbrot Computation (5)

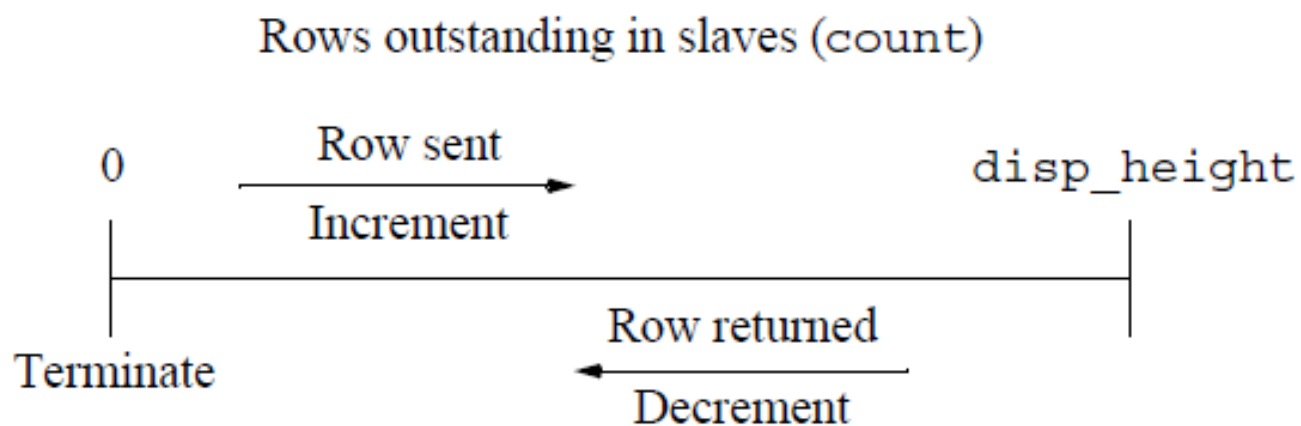


Figure 3.6 Counter termination.





# Parallelization of Mandelbrot Computation (6)

## Analysis

### Sequential

$$t_s \leq \max \times n = O(n)$$

### Parallel program

**Phase 1: Communication - Row number is sent to each slave**

$$t_{\text{comm1}} = s(t_{\text{startup}} + t_{\text{data}})$$

**Phase 2: Computation - Slaves perform their Mandelbrot computation in parallel**

$$t_{\text{comp}} \leq \frac{\max \times n}{s}$$

**Phase 3: Communication - Results passed back to master using individual sends**

$$t_{\text{comm2}} = \frac{n}{s}(t_{\text{startup}} + t_{\text{data}})$$

### Overall

$$t_p \leq \frac{\max \times n}{s} + \left(\frac{n}{s} + s\right)(t_{\text{startup}} + t_{\text{data}})$$





# Monte Carlo Methods (1)

Basis of Monte Carlo methods is the use of random selections in calculations.

## Example - To calculate $\pi$

A circle is formed within a square. Circle has unit radius so that square has sides  $2 \times 2$ .




Figure 3.7 Computing  $\pi$  by a Monte Carlo method.





## Monte Carlo Methods (2)

The ratio of the area of the circle to the square is given by

$$\frac{\text{Area of circle}}{\text{Area of square}} = \frac{\pi(1)^2}{2 \times 2} = \frac{\pi}{4}$$

Points within the square are chosen randomly and a score is kept of how many points happen to lie within the circle.

The fraction of points within the circle will be  $\pi/4$ , given a sufficient number of randomly selected samples.





# Monte Carlo Methods (3)

## Computing an Integral

One quadrant of the construction in Figure 3.7 can be described by the integral

$$\int_0^1 \sqrt{1-x^2} dx = \frac{\pi}{4}$$

A random pair of numbers,  $(x_r, y_r)$  would be generated, each between 0 and 1, and then

counted as in circle if  $y_r \leq \sqrt{1-x_r^2}$ ; that is,  $y_r^2 + x_r^2 \leq 1$ .




Figure 3.8 Function being integrated in computing  $\pi$  by a Monte Carlo method.







## Monte Carlo Methods (4)

### Alternative (better) Method

Use the random values of  $x$  to compute  $f(x)$  and sum the values of  $f(x)$ :

$$\text{Area} = \int_{x_1}^{x_2} f(x) dx = \lim_{N \rightarrow \infty} \frac{1}{N} \sum_{i=1}^N f(x_i)(x_2 - x_1)$$

where  $x_i$  are randomly generated values of  $x$  between  $x_1$  and  $x_2$ .





# Monte Carlo Methods (5)

## Example

Computing the integral

$$I = \int_{x_1}^{x_2} (x^2 - 3x) dx$$

## Sequential Code

```
sum = 0;
for (i = 0; i < N; i++) {           /* N random samples */
    xr = rand_v(x1, x2);             /* generate next random value */
    sum = sum + xr * xr - 3 * xr;    /* compute f(xr) */
}
area = (sum / N) * (x2 - x1);
```

The routine randv(x1, x2) returns a pseudorandom number between x1 and x2.





# Monte Carlo Methods (6)

## Parallel Implementation




Figure 3.9 Parallel Monte Carlo integration.





# Monte Carlo Methods (7)

## Pseudocode

Master

```
for(i = 0; i < N/n; i++) {  
    for (j = 0; j < n; j++)      /* n=no of random numbers for slave */  
        xr[j] = rand();          /* load numbers to be sent */  
    recv(P_ANY, req_tag, P_source); /* wait for a slave to make request */  
    send(xr, &n, P_source, compute_tag);  
}  
for(i = 0; i < slave_no; i++) { /* terminate computation */  
    recv(P_i, req_tag);  
    send(P_i, stop_tag);  
}  
sum = 0;  
reduce_add(&sum, P_group);
```





## Monte Carlo Methods (8)

Slave

```
sum = 0;
send(P_master, req_tag);
recv(xr, &n, P_master, source_tag);
while (source_tag == compute_tag) {
    for (i = 0; i < n; i++)
        sum = sum + xr[i] * xr[i] - 3 * xr[i];
    send(P_master, req_tag);
    recv(xr, &n, P_master, source_tag);
};
reduce_add(&sum, P_group);
```





# Random Number Generation

The most popular way of creating a pseudorandom number sequence:

$$x_1, x_2, x_3, \dots, x_{i-1}, x_i, x_{i+1}, \dots, x_{n-1}, x_n,$$

is by evaluating  $x_{i+1}$  from a carefully chosen function of  $x_i$ , often of the form

$$x_{i+1} = (ax_i + c) \bmod m$$

where  $a$ ,  $c$ , and  $m$  are constants chosen to create a sequence that has similar properties to truly random sequences.

- A good generator is with  $a = 16807$ ,  $m = 2^{31} - 1$ , and  $c = 0$ .
  - This generator creates a repeating sequence of  $2^{31} - 2$  different numbers





# Parallel Random Number Generation

It turns out that

$$x_{i+1} = (ax_i + c) \bmod m$$

$$x_{i+k} = (Ax_i + C) \bmod m$$

where  $A = a^k \bmod m$ ,  $C = c(a^{k-1} + a^{k-2} + \dots + a^1 + a^0) \bmod m$ , and  $k$  is a selected “jump” constant.




Figure 3.10 Parallel computation of a sequence.

